The Serialization Killer Language

Timm Felden

June 14, 2013

Abstract

This paper presents an approach to serializing objects, which is tailored for usability, performance and portability. Unlike other general serialization mechanisms, we provide explicit support for extension points in the serialized data, in order to provide a maximum of upward compatibility and extensibility.

To do (1)

Contents

1 Motivation

Many industrial and scientific projects suffer from platform or language dependent representation of their core data structures. These problems often cause software engineers to stick with outdated tools or even programming languages, thus causing a lot of frustration. This does not only increase the burden of hiring new project members, but can ultimately cause a project to die unnecessarily.

The approach presented in this paper provides means of platform and language independent specification of serializable data structures and therefore a safe way to let old tools of a tool suite talk to the knew ones, without even the need of recompiling the old ones. We set out to design a new language, because we believe, that the best language a programmer can use to write a new tool, is the language that he likes the most. We also hat the strict requirement to provide a solution that can describe an intermediate representation with stable parts that can be used for decades and unstable parts that may change on a daily basis, until it is know how the transported data has to be shaped.

In order to achieve this goal, we introduce two new concepts:

The first one is an easy to use specification language for data structures providing simple data types like integers and strings, abstract data types like sets and maps, type safe pointers, extension points and single inheritance. The specification language is modular allowing for more readable specifications.

The second one is a formalized mapping of specified types to a bitwise representation of stored objects. The mapping is very compact and therefore scalable, easy to understand and therefore easy to bind to a new language. It does encode the type system and can therefore provide a maximum of upward and downward compatibility, while maintaining type safety at the same time. It allows for a maximum of safety when it comes to manipulating data unknown to the generated interface, while maintaining high decoding and encoding speeds[1].

In contrast to other serialization formats such as Extensible Markup Language (XML), the serializable data can not be viewed or modified with a text editor. This however does not mean, that it is not human readable, because one can provide a human usable editor to edit arbitrary Serialization Killer Language (SKilL) files.

An improvement over XML is, that the reflective usage of stored data is expected to be quite rare, because the binding generator is able to generate an interface that ensures type safety of modifications and provides a nice integration into the target language. This leads to a situation, where it is possible to use files containing data of arbitrary types. If the data stored in the file is not used by a client, he does not have to pay for it with execution time or memory. It is also not required for a client to know the whole intermediate representation of a tool suite, but only the parts he is going to use in order to achieve his goals.

The expected file sizes range from 1 MiB to 2 GiB, while having virtually no relevant numerical limits in the file size[2]. Please note, that the skill file format is a lot more compact then equivalent XML files would be. It is expected, that files contain objects of hundreds of types with thousands of instances each. If a type in such a file would contain in average three pointers, the file size would still be around a mega byte, which is due to a very compact representation of stored data. This will also lead to high load and store performance, because the raw disk speed is expected to be the limiting factor.

---

[1] The serialization and deserialization operations are linear in the size of the input/output file.

[2] There are practical limits, such as Java having array lengths limited to 2^{31} or current file systems having a maximum file size limit that is roughly equivalent to the size of a file completely occupied by objects with a single field of a single byte. There will also be problems with raw I/O-Performance for very large files and an implementation of a binding generator, which can handle files not storable in the main memory is a tricky thing to do.

1.1 Scientific Contributions

This section is a very concise representation of contributions, that in part have already been mentioned above and in parts, will be mentioned much later.

The suggested serialization format and serialization language offer all of the following features at the same time:

- a small footprint and therefore high decoding speeds

- a fully reflective type encoding

- type safe storage of pointers both to known and unknown types[3]

- the specification language is modular[4] and easy to use

- no tool using a common intermediate representation has to know the complete specification. It is even possible to strip away or add individual fields of commonly used types.

- the coding is platform and language independent

- the coding offers a maximum of downward **and** upward compatibility

- a programmer is communicating through a generated interface, which allows programmers knowing nothing about skill to interact with it, ensures type safety easily. It also allows programmers to write tools in the language they know the best[5]

- stored data, that is never needed by a tool, will never be touched

Any of the arguments above have already been made in various contexts, but there is, to the best of our knowledge, no solution bringing all these demands together into a single product that does the job automatically.

To do (2)

1.2 Related Work

There are many approaches similar to ours, but most of them have a different focus. This section shall provide a concise list of related approaches. For potential users of SKilL, this might also present alternatives superior for individual use cases.

XML

XML is a file format (defined in [xml06]). The main differences are:

- + XML can be manipulated with a text editor[6].

- + It is easier to write a libXML for a new language than to write a SKilL backend[7].

- - XML is not an efficient encoding in terms of (disk-)space usage

---

[3]I.e. regular references and annotations.

[4]I.e. it can be distributed over many files.

[5]This is a problem especially in the scientific community, where many researchers work on similar problems but on completely different tools.

[6]Whereas skill requires a special editor, which will be provided by us eventually.

[7]This is only a relevant point if no bindings exist for the language that you want to use.

- XML is not type safe. This can be overcome partially by the XML Schema Definition Language (XSD).

- XML does not provide references to other objects out of the box.

- XML stores basically a tree, whereas a skill file contains an arbitrary amount of graphs of objects.

- XML is usually accessed through a libXML, whereas SKilL provides an API for each file format, thus a skill user does not require any SKilL skills. To be fair, there are some language bindings, mainly for Java, which offer this benefit for XML as well.

XML Schema definitions

The description language itself is more or less equivalent to most schema definition languages such as XSD (as described in [GSMT+08, PGM+08]). The downside is that schema definitions have to operate on XML and can not directly be used with a binary format. There is also no way to generate code for some client languages, including Ada, from a schema definition. For obvious reasons, there is no way to refer to other XML documents from a SKilL file. If this is a requirement, one might choose to stick with XML.

We feel that the SKilL specification language is easier to use and existing specifications are a lot easier to read then XSD files.

The type systems offered by SKilL and XSD are quite different, thus it might be worth a look which one better fits ones needs.

JAXP and xmlbeansxx

For Java and C++, there are code generators, which can turn a XML schema file into code, which is able to deal with an XML in a similar way, as it is proposed by this work. In case of Java this is even in the standard library. The downside is, that, to our knowledge, this is only possible for Java and C++, thus it leaves us with portability issues. A minor problem of this approach is the lack of support for comment generation and the inefficient storage of serialized data. An interesting observation is, that this approach deprives XML of its flexibility advantage over our solution.

ASN.1

Is not powerful enough to fit our purpose.

IDL

A concise description of IDL can be found in [Lam87]. It seems not to be powerful enough and is certainly outdated. It is so old, that there are no bindings for any modern language. There is also not much documentation on further research on that area, thus creating a new approach with similar goals but modern techniques is in fact an option.

The published format is stated to be ASCII ([Lam87] §2.4), which will cause similar efficiency problems as XML does, when large amounts of date are stored.

Apatche Thrift & Protobuf

Both lack subtyping. Protobuf has an overly complex notation language. Both seem to be optimized for network protocols, thus they do not have storage pools, which are the foundation of our serialization approach and an absolute requirement for some of our features, such as hints (see section 5.2).

Java Bytecode, LLVM/IR and others

Although Java Bytecode and the LLVM Intermediate Representation are hand crafted formats, they served as a guiding example in many ways.

Language Specific

Language specific is language specific and can therefore not be used to interface between subsystems written in different programming languages, without a lot of effort. Our aim is clearly a language independent and easy to use serialization format.

Language Interfaces

Language Interfaces do not permit serialization capabilities. Most language only provide interfaces for C, with varying quality and varying degree of automation. A significant problem are interfaces between languages with different memory models.

Interfacing between languages with different type systems and memory models over a common C interface can be very inefficient.

2 Syntax

We use the tokens `<id>`, `<string>`, `<int>` and `<comment>`. They equal C-style identifiers, strings, integer literals and comments respectively. We use a comment token, because we want to emit the comments in the generated code, in order to integrate nicely into the target languages documentation system.

2.1 The Grammar

The grammar of a SKilL definition file is defined as:

```
UNIT :=
  INCLUDE*
  DECLARATION*

INCLUDE :=
  ("include"|"with") <string> ";"?

DECLARATION :=
  DESCRIPTION
  <id>
  ((":"|"with"|"extends") <id>)?
  "{" FIELD* "}"
```

```
FIELD :=
  DESCRIPTION
  (CONSTANT|DATA) ";"?

DESCRIPTION :=
  (RESTRICTION|HINT)*
  <comment>?
  (RESTRICTION|HINT)*

RESTRICTION :=
  "@" <id> ("(" (R_ARG ("," R_ARG)*)? ")")? ";"?

R_ARG := ("%"|<int>|<string>)

HINT := "!" <id> ";"?

CONSTANT :=
  "const" TYPE <id> "=" <int>

DATA :=
  "auto"? TYPE <id>

TYPE :=
  ("map" MAPTYPE
  |"set" SETTYPE
  |"list" LISTTYPE
  |ARRAYTYPE)

MAPTYPE :=
  "<" GROUNDTYPE ("," GROUNDTYPE)+ ">"

SETTYPE :=
  "<" GROUNDTYPE ">"

LISTTYPE :=
  "<" GROUNDTYPE ">"

ARRAYTYPE :=
  GROUNDTYPE
  ("[" (<id>|<int>)? "]")?

GROUNDTYPE :=
  (<id>|"annotation")
```

Note: The Grammar is LL(1).[8]

   Comment: The optional ; at the end of includes or definitions are for convenience only.

---

[8]In fact it can be expressed as a single regular expression.

2.2 Reserved Words

The language itself has only the reserved words **annotation**, **auto**, **const**, **include**, **with**, **bool**, **map**, **list** and **set**.

However, it is strongly advised against using any identifiers which form reserved words in a potential target language, such as Ada, C++, C#, Java, JavaScript or Python. A complete list is given in appendix C.

2.3 Examples

Listing 1: Running Example

```
/** A source code location. */
SLoc {
  i16 line;
  i16 column;
  string path;
}

Block {
  SLoc begin;
  SLoc end;
  string image;
}

IfBlock : Block {
  Block thenBlock;
}

ITEBlock : IfBlock {
  Block elseBlock;
}
```

Includes, self references

Listing 2: Example 2a

```
with "example2b.skill"

A {
  A a;
  B b;
}
```

Listing 3: Example 2b

```
with "example2a.skill"

B {
```

```
    A a ;
}
```

Unicode

The usage of non ASCII characters is completely legal, but discouraged.

```
Listing 4: Unicode Support
/** some arguably legal unicode characters. */
ö {
   ö ∀;
   ö €;
}
```

3   Semantics

This section will describe the meaning of individual keywords.

3.1   Includes

The file referenced by the with statement is processed as well. The declarations of all files transitively reachable over `with` statements are collected, before any declaration in any file is evaluated.

3.2   `annotation`

The type has a tag and a size, which allows it to be inserted at any annotation locations. This is useful in order to provide extension points in the file format. The file will still be readable by older implementations, which are not able to map any meaningful type into the annotation. A language binding is expected to provide something like an annotation proxy, which is used to represent annotation objects. If an application tries to get the object behind the proxy for an object of an unknown type, this will inevitably result in an error or exception. Therefore language bindings shall provide means of inspecting whether or not the type of the object behind an annotation is known.

As we will see in section 6, annotations are roughly equivalent to the type definition

```
annotation {
  v64 baseTypeName;
  v64 basePoolIndex;
}
```

Of course, this is made transparent to the user and some language bindings will offer a special and type safe treatment of annotations.

An implementation may treat an annotation pointing to an object of unknown type like a null reference. This behavior is safe, because such an object can not exist in the serialized file, thus the annotation has not been updated upon removal of the complete type pool. This behavior might look rather strange at first glance but is an effect of lazy treatment of informations stored in skill files and completely safe.

3.3 Sub Types

A sub type of a user type can be declared by appending the keyword `with` and the super types name to a declaration. In order to be well-formed, the sub type relation must remain acyclic and must not contain unknown types.

3.4 `const`

A const field can be used in order to create guards or version numbers, as well as overwriting deprecated fields with e.g. zeroes. The deserialization mechanism has to report an error if a constant field has an unexpected value.

3.5 `auto`

The language binding will create a field with the given type, but the content is transparent to the serialization mechanism. This is useful if the inference of the content of a field is likely to be faster then storing it, e.g. if it can be inferred lazily.

3.6 Abstract Data Types

Abstract Data Types (ADTs) (in the sense of [MP88]) showed to be useful and to increase the usability and understandability of the resulting code and file format. The existential type characteristics can even be used by SKilL using restrictions and hints. In the file format, ADTs are represented using arrays and pairs as explained in detail in section 6.4.

The type system has a built-in notion of arrays, maps, lists and sets. Note that all of them are, from the view of serialization, equivalent to length encoded arrays. Their main purpose is to increase the usability of the generated Application Programming Interface (API).

3.7 Comments

Comments provided in the SKilL specification will be emitted into the generated code[9]. This approach enables users to get tool-tips in an IDE showing him this documentation. Future revisions are expected to expand the notion of a comment to match the of most documentation generation systems, such as doxygen [vH13] or javadoc [jav13].

4 The Type System

To do (3)

User types can be seen as nonempty tuples over all types. Built-in types can be wrapped in order to give them special semantics. E.g. a time stamp can be created by:

Listing 5: Time

```
time {
  /** seconds since 1.1.1970 0:00 UTC. */
  i64 date;
}
```

---

[9]If the target language does not allow for C-Style comments, the comments will be transformed in an appropriate way.
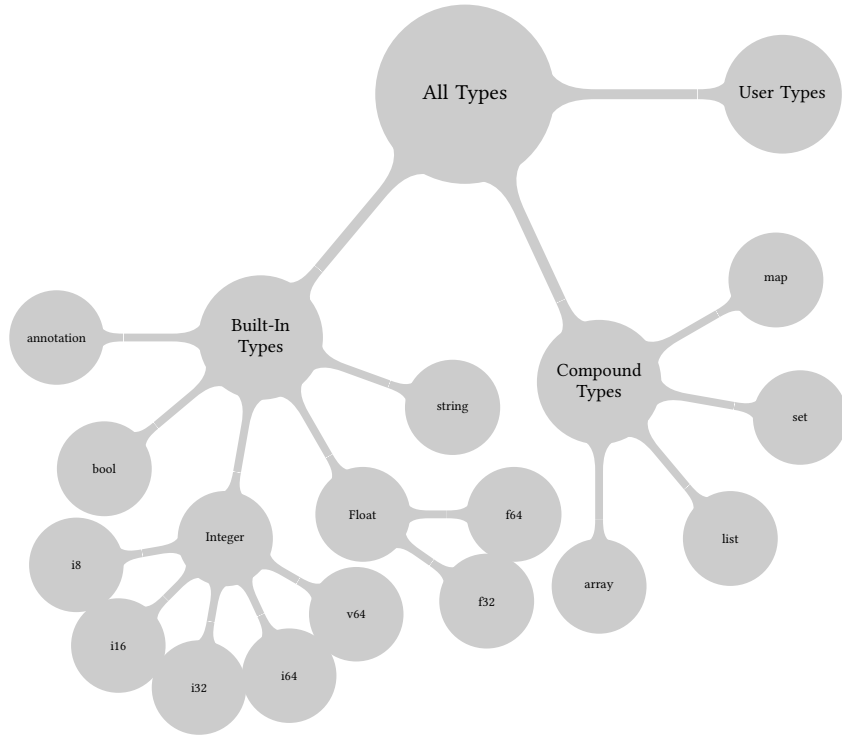
Figure 1: Layout of the Type System

Common Abbreviations

We will use some common abbreviations for sets of types in the rest of the manual:

Let …

- … $\mathcal{T}$ be the set of all types.

- … $\mathcal{U}$ be the set of all user types.

- … $\mathcal{I}$ be the set of all integer types, i.e. $\{\texttt{i8}, \texttt{i16}, \texttt{i32}, \texttt{i64}, \texttt{v64}\}$.

- … $\mathcal{B}$ be the set of all built-in types.

4.1   Legal Types

The given grammar of SKilL already ensures that intuitive usage of the language will result in legal type declarations. The remaining aspects of illegal type declarations boil down to ill-formed usage of type and field names and can be summarized as:

- Field names inside a type declaration must be unique inside the type and all its super types[10].

- The subtype relation is a partial order[11] and does not contain unknown types.

- For all fields f of dependent array type[12], the size of the array has to denote a

---

[10]The super type restriction may in fact be dropped?

[11]In fact it forms a forest.

[12]E.g. a field `t[size]  f` requires another field of integer type in the same declaration – e.g. `i8 size`

field of integer type in the very same declaration. The order of declaration is irrelevant.

- Any base type has to be known, i.e. it is either a ground type or it is a user type defined in any document transitively reachable over include commands.

4.2 Type Order

Let $<_l$ be the lexical order. We define a partial order $\leq_t$ on $\mathcal{T}$ as follows:

- $\forall t \in \mathcal{B}, s \in \mathcal{T} \setminus \mathcal{B}.t \leq_t s$

- $\forall t \in \mathcal{C}, s \in \mathcal{U}.t \leq_t s$

- $\forall s, t \in \mathcal{U}.t \leq_t s \leftarrow s <: t$[13]

- $\forall s, t \in \mathcal{U}.t \leq_t s = t \leq_l s \leftarrow \exists S \in \mathcal{U} \cup \{\bot\}.t <: S \wedge s <: S$[14]

The informal short description is, first ground types, then compound types and user types at the end, where the forest of user types maintains its structure but is order using the lexical order of type names.

Notice, that this order corresponds to an left to right order in the types overview picture.

The missing order of compound types is left away intentionally, because it allows for the exchange of some type definition after publishing a format, e.g. `t[] f` can be exchanged with `list<t> f`.

4.3 Strings

Strings are conceptually a variable length sequence of utf8-encoded unicode characters. The in memory representation will try to make use of language features such as java.lang.String or std::u16string. The serialization is described in section 6. If a language demands 0-termination in strings, the language binding will ensure this.

Strings should not contain 0 characters, because this may cause problems with languages such as C.

4.4 Compound Types

The language offers several compound types. Sets, Lists and auto sized Arrays, i.e. arrays without an explicit size, are basically views onto the same kind of serialized data, i.e. they are a length encoded list of elements of the supplied base type. Arrays are expected to have a constant size, i.e. they are not guaranteed to be resizable. Sets are not allowed to contain the same element twice. All ADTs will be mapped to their closest representation in the target language, while preserving these properties. Maps are viewed as a representation of serializable partial functions. Therefore they can contain other map types as their second type argument, which is basically an instance of currying. Maps are expected to be sparse, i.e. with less then half their fields being non default values.

---

[13]This is *super types first.*
[14]Types with the same or no supertype are order lexically.

4.5 NULL Pointer

The null pointer is serialized using the index 0. Conceptually, null pointers of different types are different. In fact if an annotation is a null pointer, it still has a type. However, this detail should not be observable in most languages.

4.6 Examples

This section will present some examples of ill-formed type declarations and brief explanations.

```
Listing 6: Legal Super Types
EncodedString : string {
  string encoding;
}
```

Error: The built-in type "string" can not be sub classed.

5 Type Annotations

5.1 Restrictions

Some invariants can be added to declarations and fields. These invariants can occur at the same place as comments, but can occur in any number. Invariants start with an @ followed by a predicate. Each predicate has to supply a default argument %, such that using only default arguments would not imply a restriction. If multiple predicates are annotated, the conjunction of them forms the invariant. The set of legal predicates is explained below.

If predicates, which are not directly applicable for compound types are used on compound types, they expand to the contents of the compound types, if applicable. Otherwise the usage of the predicate is illegal.   To do (4)

Range

Range restrictions are used to restrict integers and floats. Note that this will change the default value of the argument field to *min* if $0 \notin [min, max]$.

Applies to fields: Integer, Float.
Signature: `range(min, max)`: $\alpha \times \alpha \rightarrow bool$
Defaults: the smallest/largest value of the argument type.

```
Listing 7: Examples
natural {
  @range(0,%)
  v64 data;
}
positive {
  @range(1,%)
  v64 data;
}
nonNegativeDouble {
```

```
    @range ( 0 ,%)
    f64  data ;
}
```

NonNull

Declares that an indexed field may not be null. Note that this will take away the default value and can therefore cause compatibility problems.

    Applies to Field: Any indexed Type.
    Signature: `nonnull()`
    Defaults: none.

Listing 8: Examples

```
Node  {
    @nonnull Node [ ]  edges ;
}
```

Unique

Objects stored in a storage pool have to be distinct in their serialized form, i.e. for each pair of objects, there has to be at least one field, with a different value.

    NOTE: This can cause difficulties in combination with sub-classing, because the uniqueness property must hold even on the part restricted to the topmost class declared to be unique.

    Applies to Declarations of indexed types.
    Signature: `unique()`
    Defaults: none.

Listing 9: Examples

```
@unique  Operator  {
    string  name ;
}
@unique  Term  {
    Operator  operator ;
    Term [ ]  arguments ;
}
```

Singleton

There is at most one instance of the declaration.

    Applies to Declarations.
    Signature: `singleton()`
    Defaults: none.

Listing 10: Examples

```
@singleton  System  {  ...  }
@singleton  Data {
```

```
        /** Note: if data would not be a singleton itself, it
            is likely to violate the singleton property */
    System foo;
}
```

Ascription

A language specific type can be ascribed to a field. The type has to be compatible to the fields actual type, because the ascription will not change the ABI in any way. The first argument is the language name. The second type is generator dependent, but should be related to types as they occur in local variable or field declaration in the respective language.

Although this kind of restriction puts a heavy burden on the language generator and decreases readability a lot, it can be used to increase the usability of the generated interface a lot, because language features such es enums in Java or unions and bitfields in C++ can be used.

Applies to fields.
Signature: `as(language, type)`: $\texttt{string} \times \texttt{string} \rightarrow \{\}$
Defaults: not allowed.

Listing 11: Examples

```
System {
  /**
  The language binding makes use of an enumeration,
      which is supplied with the generated code.

  The C++ interface will use the different type using
      C−Casts to convert between the two types (which is
      completely fine if the enum uses char as a base
      type).

  The Java interface will assume the stored integer to
      be the ordinal of the enum SystemState.
  */
  @as("C++", "ccast␣SystemState")
  @as("Java", "enum␣SystemState")
  i8 state
}
```

Tree

The reference graph below created by objects of this type forms a tree. The type of the objects is irrelevant. Strings and fields with notree annotation, are not taken into account.

Applies to Declarations or Field.
Signature: `tree()`
Defaults: none.

notree

Applies to field.
    Signature: `notree()`
    Defaults: none.

Listing 12: Examples

```
Sloc {...}
@tree
SyntaktikEntity {
  /** not a tree, because several entities, might share
      them */
  @notree Sloc sloc;

  SyntaktikEntity[] children;
}
Routine {
  @notree
  Routine[] callers;
  @tree
  Routine[] dominators;
}

@tree
File {
  File[] children;
  /** several files could have the same name,
      but strings are implicitly @notree */
  string name;
  string content;
}
```

Note: In case of the File example, there is no way to violate the tree property. Note: It is legal for trees to form forests.

Constant Length Pointer

The pointer is serialized using i64 instead of v64. Can be used on regular references and annotations. This restriction makes only sense if the generated supports lazy reading of partial storage pool and if the files that have to be dealt with would not fit into the main memory of the target machine. Using this restriction will most certainly increase the file size.
    This restriction is serializable and thus, does not affect compatibility in any way.
    Applies to fields.
    Signature: constantLengthPointer()[15]
    Defaults: none.

Listing 13: Examples

---

[15] The length of the name is intended.

```
/* stored points to information may exceed the available
   main memory, thus we have to access it directly from
   disk */
PointsToTargets {
  @constantLengthPointer
  Context context;
  @constantLengthPointer
  HeapObject object;
  @constantLengthPointer
  PointsToSet targets;
}
```

5.2 Hints

Hints are annotations that start with a single ! and are followed by a hint name. Hints are used to control the behavior of the generated language binding and do not have an impact on the semantics of the stored data. Therefore they will not be stored in the reflection pool.

Access

Try to use a data structure that provides fast (random) access. E.g. an array list.

Modification

Try to use a data structure that provides fast (random) modification. E.g. a linked list[16].

Unique

Serialization shall unify objects with exactly the same serialized form. In combination with the @unique restriction, there shall at most be an error reported on deserialization.

Distributed

Use a static map instead of fields to represent fields of definitions in memory. This is usually an optimization if a definition has a lot of fields, but most use cases require only a small subset of them. Because hints do not modify the binary compatibility, some clients are likely to define the fields to be distributed or even lazy. Note that this will increase both the memory footprint and the access time for the given field and will only be a benefit for memory-cache locality reasons. The internal representation will change from `f.a`, i.e. a regular field, to `pool.a[f]`, i.e a map in the storage pool which holds the field data for each instance. Note that the presence of distributed, lazy or ignored fields will require objects to carry a pointer to their storage pool, which may eliminate the cache savings completely.

---

[16]Which has faster insert/delete operations than an array list.

Lazy

Deserialize the fields data only if it is actually used. Lazy implies distributed.

ReadOnly

The generated code is unable to modify the respective field or instances of the respective type. This options is provided to provide a consistent API while preventing from logical errors, such as modifying data from a previous stage of computation.

Ignore

The generated code is unable to access the respective field or any field of the type of the target declaration. This will lead to errors, if it is tried nonetheless. This option is provided to provide a consistent API for a combined file format, but restrict usage of certain fields, which should be transparent to the current stage of computation.

6 Serialization

This section is about representing objects as a sequence of bytes. We will call this sequence *stream*, its formal Type will be named S, the current stream will be named s. We will assume that there is an implicit conversion between fixed sized integers[17] and streams. We also make use of a stream concatenation operator $\circ : S \times S \rightarrow S$.

This section assumes, that all objects about to be serialized are already known. It further assumes, that their types and thus the values of the functions (i.e. baseType-Name, typeName, index, $[\![\_]\!]$) explained below can be easily computed.

The serialization function $[\![\_]\!]_\tau : \tau \times \mathcal{T} \rightarrow S$ will be written simply as $[\![\_]\!]$ if τ is clear form the context.

6.1 Steps of the Serialization Process

In general it is assumed that the serialization process is split into the following steps:

1. All objects to be serialized are collected. This is usually done using the transitive closure of an initial set.

2. The items are organized into their storage pools, i.e. the index function is calculated. If the state was created by deserialization and indices have changed, fields using these indices have to be updated.

3. The output stream is created as described below.

6.2 General File Layout

The file layout is optimized for lazy loading of stored data. It does also support type-safe and consistent treatment of unknown data structures. In order to achieve this, we have to store the type system used by the file together with the stored data. The

---

[17]As well as between fixed sized floating point numbers, because we define them to be IEEE-754 encoded 32-/64-bit sequences.

type system itself is using strings for its representation, thus we have motivated the following layout[18]:

```
utf8[][] stringPool

while(!EOF){
  string typeName
  string superTypeName
  option(v64 basePoolStartIndex; iff has superType)
  v64 elementCount
  [[restrictions]]
  v64 fieldCount
  foreach f in fields {
    [[f.restrictions]]
    [[f.type]]
    string f.name
    @as(f.T[elementCount])
    i8[] f.elements
  }
}
```

6.3 Storage Pools

This section contains the serialization function for an individual storage pool. We assume that storage pools are not empty. If an empty storage pool would be written to disk, it is simply skipped.[19]

Writing objects of a pool requires the following functions: $baseTypeName : \mathcal{U} \to S$, $typeName : \mathcal{U} \to S$ and $index : \mathcal{U} \cup \{\textbf{string}\} \to S$.

The basic idea behind the serialization format is to store the data grouped by type into storage pools. If objects are referred to from other objects, those references are given as an integer, which is interpreted as index into the respective storage pool. The NULL pointer is represented by the index 0.

Each pool keeps a start index, which allows for the reconstruction of the complete object. A short example (Fig. 2) shall illustrate the basic concept. It contains five types A,B,C,D and N. Each has a single field of type τ which is used to simplify the representation. The type information for the objects in the base type pool can be inferred from the data stored in the pools using the links between the base type pool and the subtype pools (The base type start index (BPSI) field of pools with a super type – shown as arrows). For the sake of readability, the name, size and count fields are omitted in the picture.

The order in which pools are serialized is currently unrestricted.

6.4 Pool Elements

In this section, we want to describe the serialization of individual fields using the function $\llbracket \_ \rrbracket_\tau$. The serialization of an objects takes places by serializing all its fields into the stream. In this section, we assume that the three functions defined in the

---

[18]The format is similar to skill itself; necessary extensions such as two dimensional variable length arrays should be self-explanatory.

[19]This has the side effect, that only type information of instantiated types are present.

19

Def.: $A\{\tau a;\}$ $N\{\tau n;\}$

$B : A\{\tau b;\}$ $C : A\{\tau c;\}$

$D : B\{\tau d;\}$

$S:$ $\cdots$ | a | a | a | a | a | a || b | b | b | b || d || c || n | $\cdots$

$\mathcal{T}:$ A B B B D C   N

BPSI:

$index:$ 1 2 3 4 5 6   2 3 4 5   5   6   1 $\cdots$

Figure 2: The serialization scheme used to store objects into pools.

last section are implicitly converted to streams using the v64 encoding. We assume further, that compound types provide a function $size : \mathcal{T} \to \mathcal{I}$, which returns the number of elements stored in a given field. Let f be a field of type t, then $[\![f]\!]$ is defined as[20]

- $\forall t \in \mathcal{U} \cup \{\textbf{string}\}.[\![f]\!]_t = \begin{cases} \texttt{0x00}, & f = \texttt{NULL} \\ index(f) & else \end{cases}$

- $[\![f]\!]_{\textbf{annotation}} = \begin{cases} \texttt{0x00 0x00}, & f = \texttt{NULL} \\ baseTypeName(f) \circ index(t) & else \end{cases}$ [21]

- $[\![\top]\!]_{\textbf{bool}} = \texttt{0xFF}$

- $[\![\bot]\!]_{\textbf{bool}} = \texttt{0x00}$

- $\forall t \in \mathcal{I} \setminus \{\textbf{v64}\}.[\![f]\!]_t = f$

- $[\![f]\!]_{\textbf{v64}} = encode(f)$ [22]

- $[\![f]\!]_{\textbf{f32}} = [\![f]\!]_{\textbf{f64}} = f$ [23]

- $\forall g \in \mathcal{B}, n \in \mathbb{N}^+ . t = g\texttt{[}n\texttt{]} \implies [\![f]\!] = [\![f_0]\!]_g \circ \cdots \circ [\![f_{n-1}]\!]_g$

- $\forall g \in \mathcal{B}, s \in \mathcal{I}, \texttt{s size}^{24} . t = g\texttt{[size]} \wedge \texttt{size} > 0 \implies [\![f]\!] = [\![f_0]\!]_g \circ \cdots \circ [\![f_{\texttt{size}-1}]\!]_g$ [25]

- $\forall g \in \mathcal{B}, n = size(f), t \in \{g\texttt{[]}, \texttt{set<}g\texttt{>}, \texttt{list<}g\texttt{>}\}.[\![f]\!] = [\![n]\!]_{v64} \circ [\![f_0]\!]_g \circ \cdots \circ [\![f_{n-1}]\!]_g$

- Maps are serialized from left to right by serializing the keyset and amending each key with the map structure which it points to. In case of Maps with two types, this is equal to a list of key value tuples. A field of type `map<T,U,V>` is serialized using a schema $[\![size(f)]\!] \circ [\![f.t_1]\!] \circ [\![size(f[t_1])]\!] \circ [\![f[t_1].u_1]\!] \circ [\![f[t_1][u_1]]\!] \circ [\![f[t_1].u_2]\!] \circ \cdots \circ [\![size(f[t_2])]\!] \circ \cdots \circ [\![f[t_n][u_m]]\!]$. Note that we treat maps like map<T, map<U,V>>.

---

[20] We will use C-Style hexadecimal integer literals for integers in streams.

[21] We do not want to use type IDs here, because we do not want to touch all annotation fields if we modify the type Pool.

[22] With encode as defined in listing 17.

[23] Assuming the float to be IEEE-754 encoded, which allows for an implicit bit-wise conversion to fixed sized integer.

[24] As stated above, size must be a field of the same declaration as f.

[25] Note that this is the only case where the encoded field does not append anything to the stream.

20

- $[\![RESTRICTION]\!] = \begin{cases} \emptyset, & id = \bot \\ [\![id]\!]_{v64}[\![arg_1]\!]_{string} \circ \cdots \circ [\![arg_n]\!]_{string} & else \end{cases}$

- $[\![t]\!]_{type} = \begin{cases} [\![id]\!]_{i8} \circ [\![val]\!]_t & id \in [0,4] \\ [\![id]\!]_{i8} & id \in [5,14] \\ 15 \circ [\![i]\!]_{v64} \circ [\![T]\!] & t = T[i] \\ 16 \circ [\![f.nameIndex]\!]_{v64} \circ [\![T]\!] & t = T[f] \\ 17 \circ [\![T]\!] & t = T[] \\ 18 \circ [\![T]\!] & t = list < T > \\ 19 \circ [\![T]\!] & t = set < T > \\ 20 \circ [\![n]\!]_{v64} \circ [\![T_1]\!] \circ \cdots \circ [\![T_n]\!] & t = map < T_i, \ldots, T_n > \\ [\![21 + reflectionPoolIndex(t)]\!]_{v64} & t \in \mathcal{U} \end{cases}$

Note that *id*s of restrictions and types a listed in appendix F.

6.5 Endianness

Files are stored in a little endian format, which is the default for common architectures.

If a client is running on a big endian machine, the endianness has to be corrected, both when reading and writing files. This can be done by changing the implementation of $[\![\_]\!]_{i*}$- and $[\![\_]\!]_{f*}$-translations.

7 Deserialization

Deserialization is mostly straight forward.

The general strategy is:

- the string pool is deserialized into an array

- the reflection pool is deserialized using the strings array

- the structure of storage pools is read, pools are created and chunks of field data are copied into memory

- required fields are parsed using the information from the reflection pool

Date Example

Let d be the deserialization function – basically the inverse function of $[\![\_]\!]$.

$d(01046461746501000200010000B010A01FFFFFFFFFFFFFFFFFF)$
$= d(01)d(046461746501000200010000B010A01FFFFFFFFFFFFFFFFFF)$
$= d(01)d(04)d(6461746501000200010000B010A01FFFFFFFFFFFFFFFFFF)$
$= d(01)d(04)d(64617465)d(01000200010000B010A01FFFFFFFFFFFFFFFFFF)$
$= string[1 : "date"]d(01000200010000B010A01FFFFFFFFFFFFFFFFFF)$
$= string[1 : "date"]d(01)d(00)d(02)d(00)d(01)d(0000B010A01FFFFFFFFFFFFFFFFFF)$
$= string[1 : "date"]date[T : date[\_]1 : \_2 : \_d(0000B010A01FFFFFFFFFFFFFFFFFF)$
$= string[1 : "date"]date[T : date[d(0000B01)]1 : \_2 : \_d(0A01FFFFFFFFFFFFFFFFFF)$
$= string[1 : "date"]date[T : date[v64date]1 : \_2 : \_d(0A01FFFFFFFFFFFFFFFFFF)$
$= string[1 : "date"]date[T : date[v64date]1 : \_2 : \_d(01FFFFFFFFFFFFFFFFFF)]$
$= string[1 : "date"]date[T : date[v64date]1 : [1]2 : [-1]]$

8 In Memory Representation

This section is to describe the API provided to the programmer and the representation of objects inside memory. Both are not mandatory and can depend on the target language and implementation.

8.1 API

The generated API has to be designed in a fashion that integrates nicely with the languages programming paradigms. E.g. in Java it would be most useful to create a state object, which holds state of a bunch of serializable data and provides iterators over existing objects, as well as factory methods and methods to remove objects form the state object. The serialized types can be represented by interfaces providing getters, setters, using hidden implementations, only known to the state object.

talk about the generated API and its features, like iterators, factories, access to singletons and stuff.

Examples

Nice example in C++:

Listing 14: C++ Examples

```cpp
#include <stdint.h>
#include <string>
[...some other bouilerplate includes...]
struct SLoc {
  uint16_t line;
  uint16_t column;
  std::string* path;
};
struct Block {
  std::string* tag;
  SLoc* begin;
  SLoc* end;
  std::string* image;
};
struct IfBlock : public Block {
  Block thenBlock;
};
struct ITEBlock : public IfBlock {
  Block elseBlock;
};
[...
  plus some boilerplate code for visitors, iostreams etc.
...]
```

Listing 15: Java Examples

```java
class SLoc {
```

```
  public short line;
  public short column;
  public String path;
}
class Block {
  final public String tag() {
    return this.getClass().getName();
  }
  public SLoc begin;
  public SLoc end;
  public String image:
}
class IfBlock extends Block {
  public Block thenBlock;
}
class ITEBlock extends IfBlock {
  public Block elseBlock;
}
[...some read and write code, plus some visitors ...]
```

Listing 16: LaTeX Examples

```
$(line, column, path) \in SLoc
  \subseteq \mathbb{Z} \times \mathbb{Z} \times string$

$(begin, end, image) \in Block
  \subseteq SLoc \times SLoc \times string$

$(super, thenBlock) \in IfBlock
  \subseteq Block \times Block$

$(super, elseBlock) \in ITEBlock
  \subseteq IfBlock \times Block$
```

Which looks like:

$(line, column, path) \in SLoc \subseteq \mathbb{Z} \times \mathbb{Z} \times string$

$(begin, end, image) \in Block \subseteq SLoc \times SLoc \times string$

$(super, thenBlock) \in IfBlock \subseteq Block \times Block$

$(super, elseBlock) \in ITEBlock \subseteq IfBlock \times Block$

Note: The incentive of the LaTeX-output is to provide a mechanism for users to formalize their file format using mechanisms, that are or can not be available as a specification language. E.g. the sentence "The path of a SLoc points to a valid file on the file system and the line and column form a valid location inside that file." can not be verified in a static manner. This is because the correctness of the property depends not only on the content to be verified, but on the verifying environment as well.

8.2 Representation of Objects

The combination of laziness and consistency has the effect, that representation of objects inside memory is rather difficult. This section describes data structures and

algorithms which basically do the job. In this section, we assume that all fields are present as arrays of bytes. We will describe the effects of parsing fields in an unmodified state, in a modified state, how to modify a state and finally how to write a state back to disk.

8.2.1 Proposed Data Structures

A state has to contain at least these informations:

- an array of strings

- the type information

- storage pools

A storage pool has to hold the images of fields, which are not yet parsed.

Objects are required to have an ID field, which corresponds to the ID of the deserialized(!) state. This field is required in order to map the lazy fields to the correct objects. It can also be reused in the serialization phase to assign unique IDs, which will be used instead of pointers.

Objects of types with eager fields should have the respective fields. E.g. the declaration T {t a; !lazy t b;} should be represented by an object

```
InternalTObject {
  long ID;
  t a;
  /* getA, setA... */

  StoragePool tPool;
  /* getB, setB... */
}
```

Note that the pointer to the enclosing storage pool is required for the correct treatment of lazy and distributed fields. This is because the pool holds the field data.

Now that we have a representation for objects, we still have to store objects across storage pools. The possibility of inheritance requires us to store objects in multiple pools at the same time. We propose to store super and sub type information inside pools as links to the pools of the respective types. The objects should be stored as a "double linked array list", which is basically a linked list containing array lists:

Each pool stores the objects, with the static type of the pool in an array list. The pools of subtypes are stored in a linked list. Now an iterator over all elements of a type uses an iterator over all elements of a pool in combination with iterators over the pool and all its sub pools. Therefore creating and deleting objects is an amortized O(1) operation and it is guaranteed to maintain the semantic structure of the file, if IDs are not updated in phases other than reading and writing a file.

Note that, in the presence of distributed fields, the runtime complexity of these operations will change. It is expected to reduce to O(log(n)) where n is the largest number of objects with a common distributed field.

8.2.2 Reading Unmodified Data

Reading unmodified data is basically done by creating objects with ascending IDs and adding all eagerly processed fields to them. Pointer resolution in an unmodified state

is an O(t) operation, where t is the size of the type hierarchy below the static type of the pointer.

During the reconstruction of the initial dataset, an array in the base pool may be used to reduce the cost to O(1)[26]. However, this helper array has to be dropped, as soon as the base pool is modified.

8.2.3 Modifying Data

The only legal way of modifying data is to access it through the generated API, which provides iterators, a type safe facade, factories and means of removing objects from states for each known type. A modification is any operation that will invalidate any existing object ID, i.e. deleting objects or inserting objects into nonempty storage pools. Adding objects to empty storage pools does not count as a modification in the sense of a modified state, because it is not possible, that a pointer to such an object lures in an yet untreated field.

8.2.4 Reading Data in an Modified State

Reading Data in a modified state, is very similar to reading data in the unmodified state. Except that resolution of stored pointers can no longer rely on the trick that the ID of an object is also the index into the base type pool. There is a solution for this problem using O(t + log(n)), where n is the number of instances with the same static type. However, the straightforward implementation is O(t+n).

With this difference in mind, we strongly recommend adding a dirty flag to each storage pool which traces modifications. This will effectively eliminate the additional cost, because transformation of stored state is expected to be monotonic growing in a way that each step adds instances of a type, which was not present before and does not change old data.

8.2.5 Writing Data

Data which is lazy and modified can not be wrote to disk. Therefore any data which can potentially refer to modified data has to be evaluated. After evaluating the respective data, serialization is straightforward. The evaluation of lazy data referring to potentially modified data can be done in $O(out)$, where out is the size of the output file. Writing the output is also in $O(out)$, thus it is not per se a problem.

8.2.6 Final Thoughts on Runtime Complexity

Although the last sections read a lot like accessing serializable data being unnecessary expensive, this is in fact not the case.

Reading data without modifying it is $O(in)$, even if only a part of the data is read. This is mainly caused by the requirement of being able to process unknown data correctly. The actual cost should be limited by the cost of sequentially reading the input file from disk.

Reading data modifying it and writing it back is $O(in+m+out)$, which is not surprising at all, because one has to pay for reading the initial file, writing the complete output file and the modifications.

---

[26]The creation time for the array can be paid for during the creation of the base pool.

Only the usage of a lot of lazy or distributed data is expensive. It is generally advised against using the lazy attribute if a field is read for sure during the lifetime of a serializable state.

9 Future Work

XML output mit XML Schema.

Das neue Serialisierungsschema erlaubt es einen Viewer zu bauen, der Definition+Datei anzeigen kann. (Die future work ist hier der viewer)

Integration der Definition in die Serialisierte Form, damit man die Daten generisch prüfen und anzeigen kann. Hier braucht man noch ein gutes encoding, weil man sonst zu viel platz verbraucht.

Abuse annotations for type-safe unions. The type system does not allow for unrestricted unions or intersection types. The former violate serialization invariants, the latter would either have no instances or be equal to an already existing (super) type.

State somewhere, that a major advantage over XML is, that one is not required to link against a overly general implementation, which is nice if one is only interested in a very specific format.

A notion of *first class strings* can be used to separate the string pool into one pool at the beginning of the file which contains all the strings, that have to be used in order to understand the contents of the file and a second part of the pool, which can be skipped. This should give significant performance improvements if files with lots of unused strings are processed.

True comments with # ... \n?

Alternativ kann man die reflection data aus dem ReflectionPool als Eingabesprache für einen Generator benutzen.

Fun fact: Garbage Collection for serializable objects comes for free, if objects are always held in storage pools.

Add generic/template import statements, which allow to import files together with a substitution. That way one can create more complex ADTs such as B-Trees. This feature is not a priority and only useful for large files and projects and requires ascription to yield the desired effect automatically.

Add *add* and *subtract* declarations to the file format, starting with ++ or −, which allow taking away or adding fields to types. If this is done, a usability evaluation is necessary!

To do…

☐ 1 (p. 1): namensäquivalenz muss immernoch mit lowercase arbeiten, wie es u.a. in Ada ist; die Konvention ist CamelCase und backends sind angehalten CamelCase in die entsprechende konvention der ausgabesprache zu übertragen, z.B. indem man in ada alles in uppercase schreibt und vor den inneren großbuchstaben einen _ einfügt, falls es nur einer ist

☐ 2 (p. 4): cite, cite, cite; llvm tutorial "llvm and perl", platform independent javabytecode, IDL, ada serialization, custom binary formats

☐ 3 (p. 10): ein paar einleitende worte

☐ 4 (p. 13): hier muss man zwischen serialisierbaren und nicht serialisierbaren restrictions unterscheiden; serialisierbar sind alle restrictions, die auch auswirkungen auf die potentiell gespeicherten daten haben, wie range und nonnull

☐ 5 (p. 22): Hier muss man noch etwas darüber sagen, was man macht, wenn die in memory representation nicht mit der repräsentation in datein übereinstimmt. Hier bieten sich die beiden formen invariante typen und größere typen an. Das format darf sich im contravianten fall aber nicht ändern, weil man sonst die aufwärtskompatibilität gefährden würde. Man kann noch über einen modus *enforced downcasts* nachdenken, der allerdings unweigerlich fehler verursachen wird, wenn downcasts bei classen auftreten, die inkompatibel sind. Außerdem werden defaultwerte erzeugt und die semantik von integer feldern ist gefährdet.

Part I
Appendix

A   Variable Length Coding

Size and Length information is stored as variable length coded 64 bit unsigned integers (aka C's `uint64_t`). The basic idea is to use up to 9 bytes, where any byte starts with a 1 iff there is a consecutive byte. This leaves a payload of 7 bit for the first 8 bytes and 8 bits of payload for the ninth byte. This is very similar to the famous utf8 encoding and is motivated, as it is the case with utf8, by the assumption, that smaller numbers are a lot more likely. It has the nice property, that there are virtually no numerical size limitations. The following small C++ functions will illustrate the algorithm:

Listing 17: Variable Length Encoding

```
uint8_t* encode(uint64_t v){
  // calculate effective size
  int size = 0;
  {
    auto q = v;
    while(q){
      q >>= 7;
      size++;
    }
  }
  if(!size){
    auto rval = new uint8_t[1];
    rval[0]=0;
    return rval;
  } else if(10==size)
    size = 9;

  // split
  auto rval = new uint8_t[size];
  int count=0;
  for(;count<8&&count<size-1;count++){
    rval[count] = v >> (7*count);
    rval[count] |= 0x80;
  }
  rval[count] = v >> (7*count);
  return rval;
}
```

```
uint64_t decode(uint8_t* p){
  int count = 0;
  uint64_t rval = 0;
  register uint64_t r;
  for(;count<8 && (*p)&0x80; count++, p++){
    r = p[0];
    rval |= (r&0x7f)<<(7*count);
  }
  r = p[0];
  rval |= (8==count?r:(r&0x7f))<<(7*count);
  return rval;
}
```

B Error Reporting

This section describes some errors regarding ill-formatted files, which must be detected and reported. The order is based on the expected order of checking for the described error. The described errors are expected to be the result of file corruption, format change or bugs in a language binding.

Deserialization

- If `EOF` is encountered unexpectedly, an error must be reported before producing any observable result.

- If an index into a pool is invalid[27], an error must be reported.

- If the deserialization of a storage pool does not consume exactly the `sizeBytes` in its header, an error must be reported. Note: This is a strong indicator for a format change.

- If the serialized type order of storage pools does not match the expected type order, an error must be reported.

- If the serialized type information contains cycles, an error must be reported, which contains at least all type names in the detected cycle and the base type, if one can be determined.

- If a storage pools contains elements which, based on their location in the base pool, should be subtypes of some kind, but have no respective sub type storage pool, an error must be reported with at least, the base type name, the most exact known type name and the adjacent base type names.

- All known constant fields have to be checked before producing any observable result. If some constant value differs from the expected value, an error must be reported, which contains at least the type, the field type and name, the base-PoolIndex, the index inside the types pool, the expected value and the actual value.

- If a serialized value violates a restriction or the invariant of a type,[28] an error must be reported as soon as this fact can be observed. It is explicitly not required to check all serialized data for this property.

C   Reserved Words

This section contains a table of words which must not be used as field names, because they are keywords in some languages. The usage of skill keywords will result in a direct error, whereas the usage of a word listed below will result in a warning, because the identifier will be escaped in the target language binding.

if	**then**	**else**	**begin**	**end**
struct	**class**	**public**	**protected**	**private**
$\Rightarrow$		…		

D   Core Language

The core language is a subset of the full language which must be supported by any generator, which is called skill core language generator. Features included in the core language are:

- Integer types `i8` to `i64` and `v64`

---

[27]because it is larger then the last string in the pool

[28]Including sets containing multiple similar elements.

- `string`, `bool` and `annotation`

- Compound types

- User Types with sub-typing

- `const` and `auto` fields.

- Reflection.

Thus the remaining parts required for full skill support are:

- Floats

- Restrictions

- Hints

- Language dependent treatment of comments, e.g. integration into doxygen or javadoc.[29]

- Name mangling to allow for usage of language keywords or illegal characters (unicode) in specification files, without making a language binding impossible.

E   Numerical Limits

In order to keep serialized data platform independent, one has to respect the numerical limits of the various target platforms. For instance, the Java Virtual Machine will not allow arrays with a size larger then 2^31 minus some elements. Therefore we establish the following rule:

(De-)serialization of a file with an array of more then 2^30 elements or a type with more then 2^30 instances may fail due to numerical limits of the target platform.

F   Numerical Constants

This section will list the translation of type IDs(as required in section 6.4) and restriction IDs (see section 5.1 and 6.4). Restrictions with undefined IDs will not be serialized.

---

[29]This may even require a language extension providing tags inside comments which are translated into tags of the respective documentation framework.

Type Name	Value
const i8	0
const i16	1
const i32	2
const i64	3
const v64	4
annotation	5
bool	6
i8	7
i16	8
i32	9
i64	10
v64	11
f32	12
f64	13
string	14
T[i]	15
T[f]	16
T[]	17
list<T>	18
set<T>	19
map<T_1, ..., T_n>	20
T	21 + $index_T$

(a) Type IDs

Restriction Name	Value
range	0
nonnull	1
unique	2
singleton	3
...	...

(b) Restriction IDs

Glossary

base type The root of a type tree, i.e. the farthest type reach able over the super type relation.. 21, 38

built-in type Any predefined type, that is not a compound type, i.e. annotations, booleans, integers, floats and strings.. 11, 38

ground type Any type, that is not a compound type, i.e. the union of user defined types and built-in types. 38

sub type If a user type A extends a type B, A is called the sub type (of B).. 10, 38

super type If a user type A extends a type B, B is called the super type (of A).. 10, 38

unknown type We will call a type *unknown*, if there is no visible declaration of the type. Such types must not occur in a declaration file, but they can be encountered in the serialization or deserialization process.. 10, 38

user type Any type, that is defined by the user using a type declaration. 10, 11, 38

visible declaration We will call a type declaration *visible*, if it is defined in the local file, or in any file transitively reachable over include directives.. 38

Acronyms

ABI Application Binary Interface. 38

ADT Abstract Data Type. 11, 38

API Application Programming Interface. 11, 24, 38

SKilL Serialization Killer Language. 3–5, 7, 12, 38

XML Extensible Markup Language. 3, 5, 38

XSD XML Schema Definition Language. 5, 38

References

[GSMT⁺08] Shudi Gao, C. Michael Sperberg-McQueen, Henry S. Thompson, Noah Mendelsohn, David Beech, and Murray Maloney. W3c xml schema definition language (xsd) 1.1 part 1: Structures. World Wide Web Consortium, Working Draft WD-xmlschema11-1-20080620, June 2008.

[jav13] Javadoc technology, 2013.

[Lam87] David Alex Lamb. Idl: Sharing intermediate representations. *ACM Trans. Program. Lang. Syst.*, 9(3):297–318, 1987.

[MP88] John C. Mitchell and Gordon D. Plotkin. Abstract types have existential type. *ACM Trans. Program. Lang. Syst.*, 10(3):470–502, July 1988.

[PGM⁺08] David Peterson, Shudi Gao, Ashok Malhotra, C. Michael Sperberg-McQueen, and Henry S. Thompson. W3c xml schema definition language (xsd) 1.1 part 2: Datatypes. World Wide Web Consortium, Working Draft WD-xmlschema11-2-20080620, June 2008.

[vH13] Dimitri van Heesch. *Doxygen User Manual.* http://www.stack.nl/ dimitri/doxygen/manual/, 2013.

[xml06] Extensible markup language (xml) 1.1 (second edition). W3c recommendation, W3C - World Wide Web Consortium, September 2006.